

ASL to Audio

Autobot: Curtis Lin, Satya Thiruvallur, Param Viswanathan, Hailey Wu
W251 Project, 2020 Spring

Agenda

Background + Dataset

Recorded Demo

Pipeline

Cloud

Edge

Live Demo

Next Step


Background + Dataset

ASL Challenge: Communication requires both parties (or multiple participants) to understand and use sign languages as the prerequisite.

Project Goal: The Autobot team wants to overcome challenges for the deaf-mute community, by converting the sign languages directly to audio.

Dataset Options: word-level vs. alphabet-level


Word level



Letter level



Recorded Demo



Dataset selection and implementation plan

Cloud:

Dataset: ASL Alphabet
(1.2 GB; 87000 images;
29 classes)



Cloud VM
(Model training: VGG16, Resnet 50,
Custom Vision, AutoML)

Edge (Jetson TX2):

USB
camera



Inference on Edge
(pretrained model:
Keras, OpenCV)

Pipelines:

1. Real-time video capturing
2. Hand sign image collection
3. Sign prediction
4. Text printing
5. Save the audio of a sentence

CLOUD

Cloud – Model Architecture



Approach	Strengths	Weakness
Build from scratch	<ul style="list-style-type: none">❖ Most flexible architecture❖ Freedom to tune parameters❖ Shot at greatness (<i>SatyaNet</i>)	<ul style="list-style-type: none">❖ Longer development time❖ Computational resource limitation❖ Reinventing the wheel
Bottleneck features (SOTA)	<ul style="list-style-type: none">❖ Leverage features trained on ImageNet❖ Performant❖ Faster dev time	<ul style="list-style-type: none">❖ Higher Accuracy but can go higher❖ Features may not be relevant
Fine-tune top layer (SOTA)	<ul style="list-style-type: none">❖ Highest Model Accuracy❖ Fastest dev time❖ Incentivized by HW9 losses	<ul style="list-style-type: none">❖ Incompatible pre-trained weights❖ Overfitting by powerful networks

VGG-16 Transfer Learning



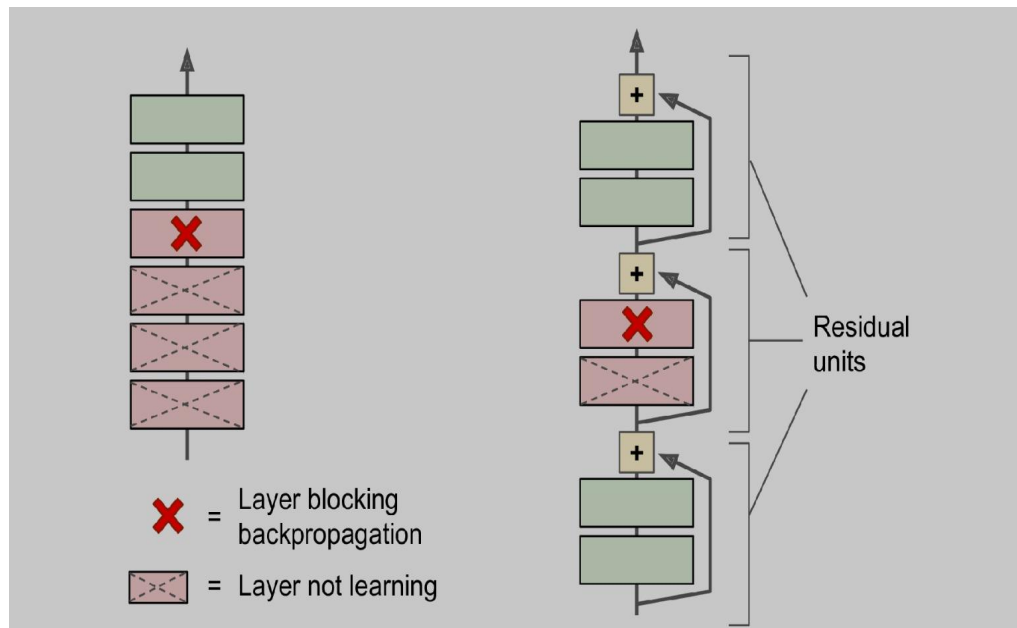
- ❖ Large kernel sized filters replaced by multiple small 3X3 filters
- ❖ Trained with 87000 50X50 sized ASL images. Labels were One Hot encoded
- ❖ Dimension ordered pre-trained VGG-16 weights used without top layer
- ❖ Categorical Cross entropy loss , Adam optimizer
- ❖ Total params: 14, 729, 565 Trainable Prams: 14, 877

Key Metrics



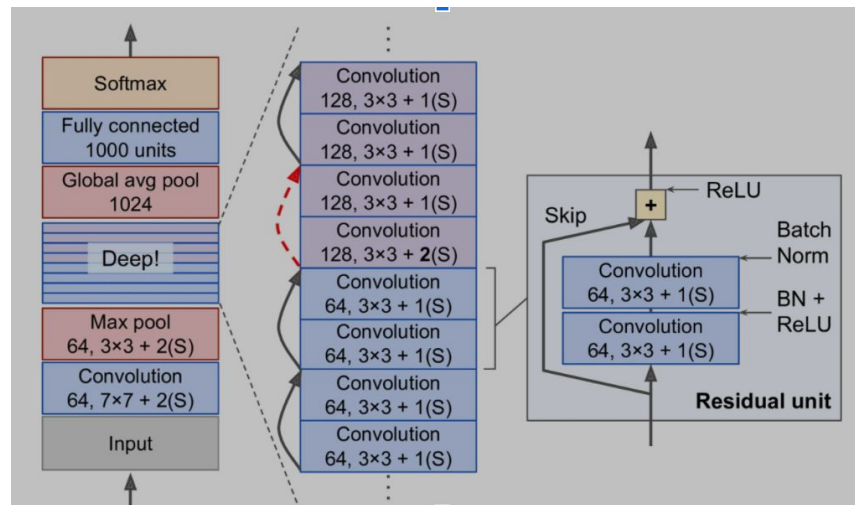
Cloud – Model ResNet50

- Residual Network with 50 layers
- Deep network that uses skip connection strategy
- Stack of residual units (RUs) where each unit is a small neural network with a skip connection



Cloud – Model ResNet50

- Deep stack of simple residual units
- Each unit is composed of 2 convolutional layers (no pooling layer)
- Batch Normalization and ReLU activation
- 3x3 Kernels
- Preserves spatial dimensions (stride=1, “same” padding)



Cloud – Model ResNet50

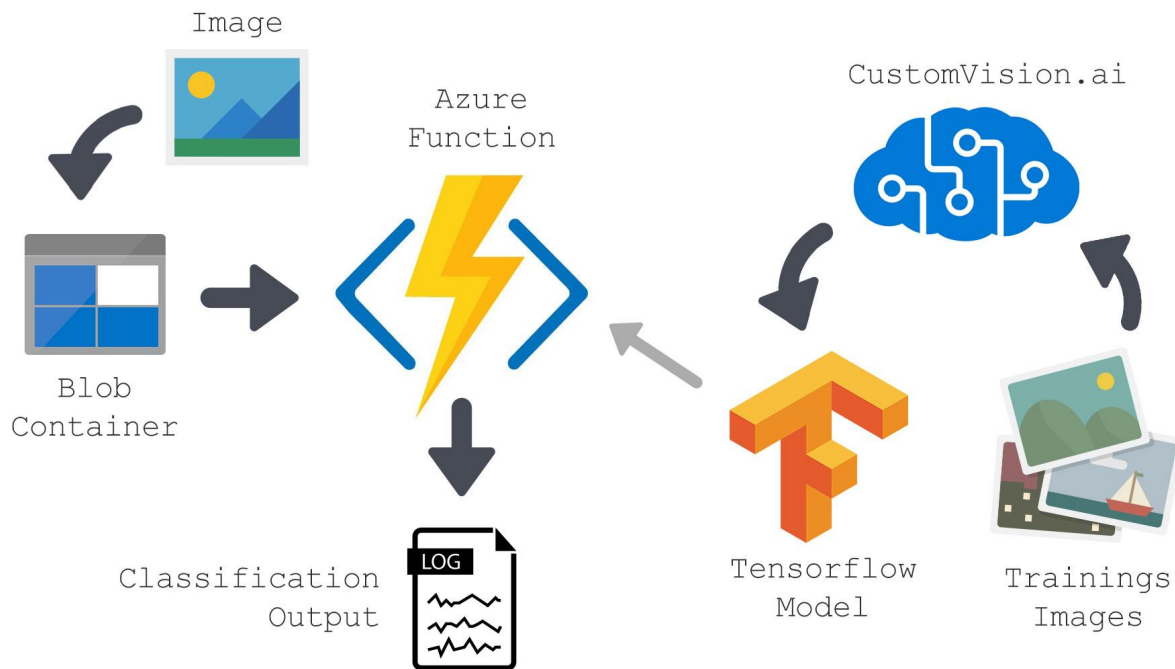

Training Strategy

- Image was pre-processed to a size of 224 x 224 and augmented
- Top layer of the network was replaced with our Pool and Dense layers
- Two runs of 10 epochs each were executed
- First run
 - Weights of pretrained layers were frozen
 - Learning rate = 0.2 and decay = 0.01
 - Accuracy = 65%
- Second run
 - Weights of pretrained layers were not frozen
 - Learning rate = 0.01 and decay = 0.001
 - Accuracy = 98%


Cloud – Model ResNet50

```
Epoch 1/10
2446/2446 [=====] - 458s 187ms/step - loss: 0.2329 - accuracy: 0.9489 - val_loss: 0.0624 - val_accuracy: 0.9778
Epoch 2/10
2446/2446 [=====] - 455s 186ms/step - loss: 0.0033 - accuracy: 0.9997 - val_loss: 0.0535 - val_accuracy: 0.9837
Epoch 3/10
2446/2446 [=====] - 455s 186ms/step - loss: 0.0018 - accuracy: 0.9999 - val_loss: 0.0586 - val_accuracy: 0.9823
Epoch 4/10
2446/2446 [=====] - 455s 186ms/step - loss: 0.0015 - accuracy: 0.9999 - val_loss: 0.0518 - val_accuracy: 0.9844
Epoch 5/10
2446/2446 [=====] - 455s 186ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.0571 - val_accuracy: 0.9828
Epoch 6/10
2446/2446 [=====] - 455s 186ms/step - loss: 0.0011 - accuracy: 0.9999 - val_loss: 0.0548 - val_accuracy: 0.9832
Epoch 7/10
2446/2446 [=====] - 454s 186ms/step - loss: 9.1674e-04 - accuracy: 0.9999 - val_loss: 0.0660 - val_accuracy: 0.9805
Epoch 8/10
2446/2446 [=====] - 455s 186ms/step - loss: 8.7085e-04 - accuracy: 0.9999 - val_loss: 0.0522 - val_accuracy: 0.9851
Epoch 9/10
2446/2446 [=====] - 454s 186ms/step - loss: 8.5705e-04 - accuracy: 1.0000 - val_loss: 0.0553 - val_accuracy: 0.9829
Epoch 10/10
2446/2446 [=====] - 454s 186ms/step - loss: 7.6152e-04 - accuracy: 0.9999 - val_loss: 0.0534 - val_accuracy: 0.9838
```

Cloud – Custom Vision (Azure)



Cloud – Custom Vision (Azure)



ASL Recognition

Training Images Performance Predictions Train Quick Test

Filter

Iteration

Workspace

Tags

Tagged Untagged

Showing: all tagged images

Search For Tags:

- ☐ C 3070
- ☐ del 3000
- ☐ M 3068
- ☐ N 3070
- ☐ O 3070
- ☐ P 3070
- ☐ space 3000
- ☐ U 3070
- ☐ V 3070
- ☐ W 3070


Show 19 more

Add images Delete Tag images Select all


1 2 3 4 5

Get started


Cloud – AutoML (GCP)



How AutoML works



Cloud – AutoML (GCP)



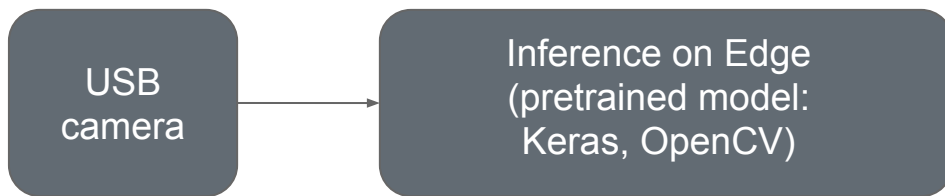
Set up process

- Copy a set of images into Google Cloud Storage.
- Create a CSV listing the images and their labels.
- Use AutoML Vision to create your dataset, train a custom AutoML Vision Edge model, and make a prediction.
- Export and deploy the AutoML Vision Edge model.

EDGE


Edge-pipelines

Edge (Jetson TX2):




The pipeline has five major steps:

- Read the real-time video input
- Cut out the hands portion from the frame
- Conduct inference on the hands picture
- Return the text
- Return the audio of a sentence




Interface—Windows of image capturing



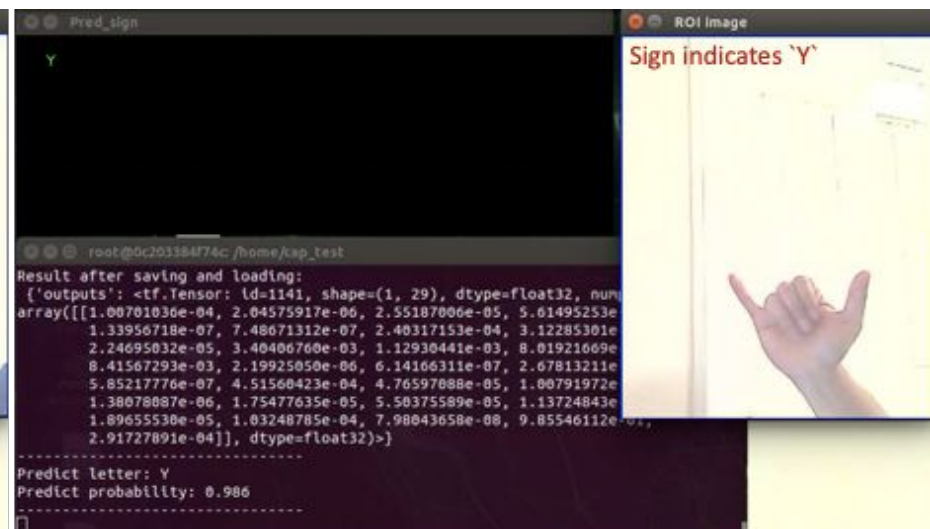
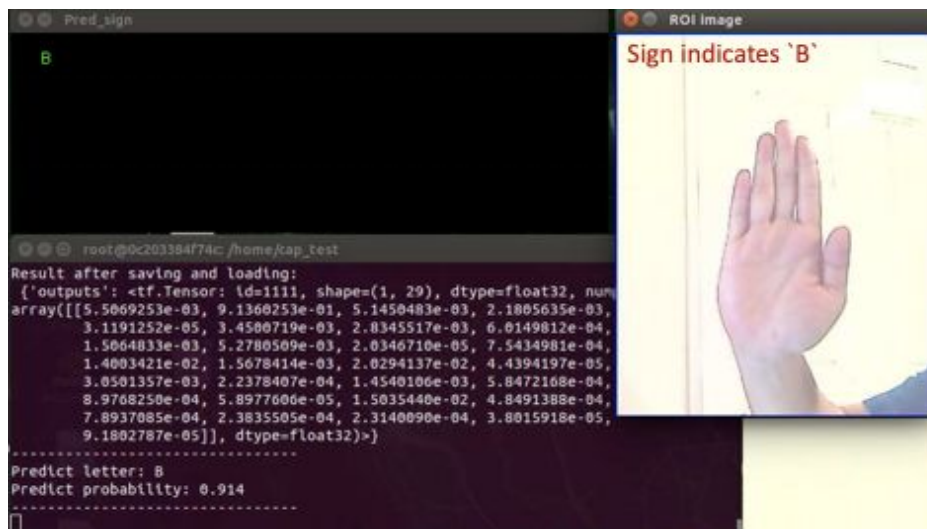
Interface-Prediction results

```
root@0c203384f74c: /home/cap_test
Result after saving and loading:
{'outputs': <tf.Tensor: id=2861, shape=(1, 29), dtype=float32, numpy=
array([[3.6763038e-05, 8.1548101e-04, 1.1495011e-02, 8.7623578e-01,
1.9593233e-05, 1.1486478e-05, 6.7616485e-02, 1.3816097e-05,
8.8108945e-06, 2.7282885e-03, 4.8795091e-06, 9.8892802e-04,
4.9960795e-03, 2.0433907e-05, 1.0951474e-04, 1.4912941e-05,
2.7519494e-02, 1.0796617e-03, 1.9500379e-05, 6.8796013e-05,
3.4988641e-06, 8.9157716e-04, 2.9054610e-04, 9.9610743e-06,
2.1294598e-03, 1.6401726e-04, 2.6805883e-03, 8.8748948e-06,
1.7797596e-05]], dtype=float32)>}]
-----
Predict letter: D
Predict probability: 0.876
-----
```

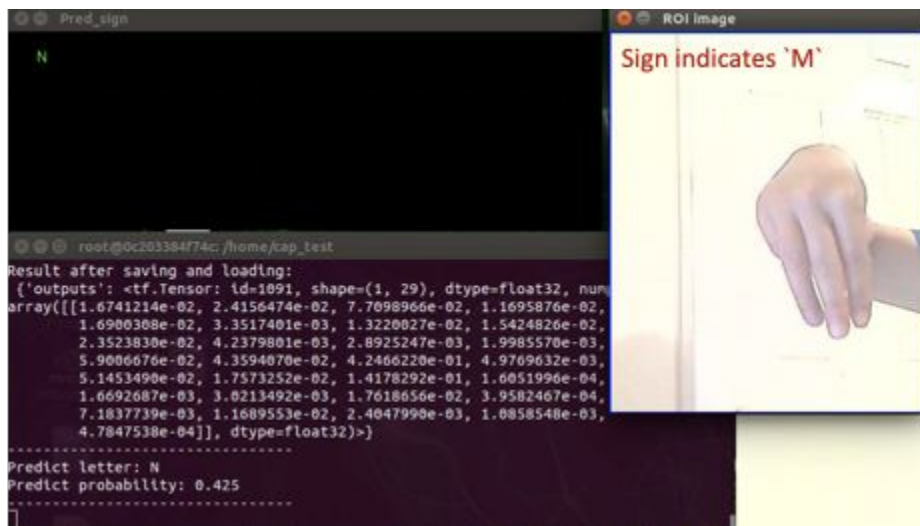
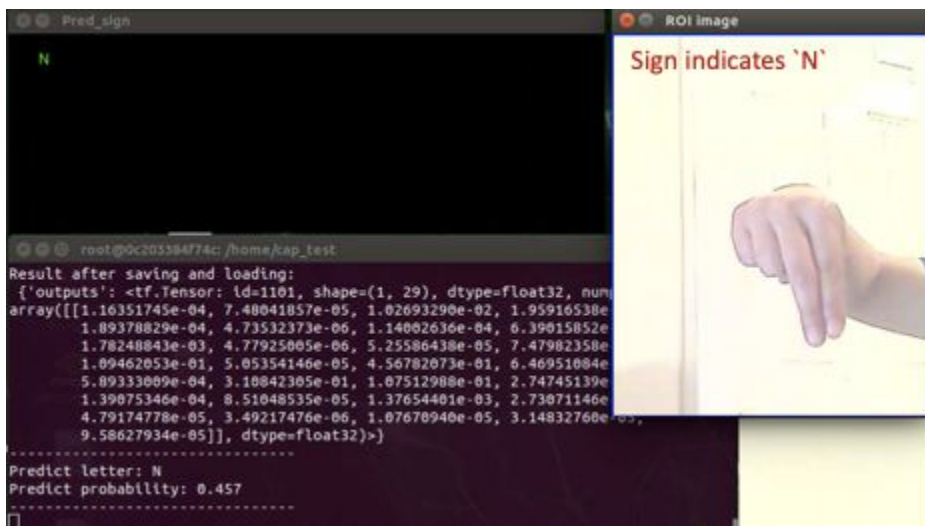
Interface-Sign language to text and audio



Prediction



Issue of misclassification




Live Demo

Next Steps–Improvement of prediction


Image preprocessing

Image augmentation

Photo auto-tuning



Palm Detection



Next Steps–Expansion with word–level signs

Hungry



Rice



Next Steps–Mobile Apps

Mobile App



Thank You

Reference

- <https://arxiv.org/abs/1910.11006>
- <https://arxiv.org/abs/1409.1556>
- <https://www.geeksforgeeks.org/convert-text-speech-python/>
- https://github.com/athena15/project_kojak/blob/master/real_time_gesture_detection.py
- <https://ai.googleblog.com/2019/08/on-device-real-time-hand-tracking-with.html>
- https://docs.opencv.org/3.4/d3/dc1/tutorial_basic_linear_transform.html
- <https://blog.metaflow.fr/tensorflow-how-to-freeze-a-model-and-serve-it-with-a-python-api-d4f3596b3adc>
- https://www.tensorflow.org/guide/saved_model
- <https://towardsdatascience.com/a-comprehensive-hands-on-guide-to-transfer-learning-with-real-world-applications-in-deep-learning-212bf3b2f27a>
- <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>
- <https://medium.com/@prasadpal107/saving-freezing-optimizing-for-inference-restoring-of-tensorflow-models-b4146deb21b5>
- <https://github.com/victordibia/handtracking>
- https://docs.opencv.org/3.4/d3/dc1/tutorial_basic_linear_transform.html
- <https://leimao.github.io/blog/Save-Load-Inference-From-TF-Frozen-Graph/>
- <https://github.com/fchollet/deep-learning-models/releases>
- https://github.com/pipidog/keras_to_tensorflow#load-a-pb-file-by-tensorflow